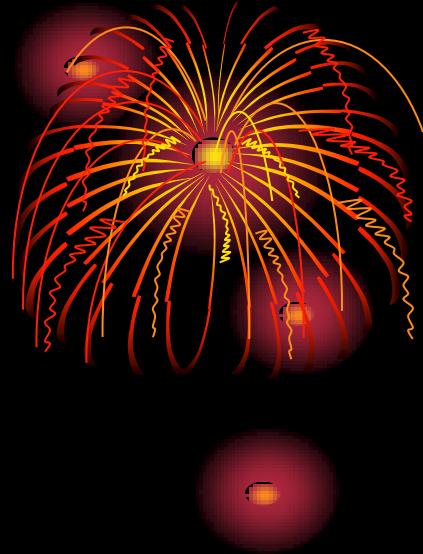


Parallel Programming

Parallel algorithms Sorting

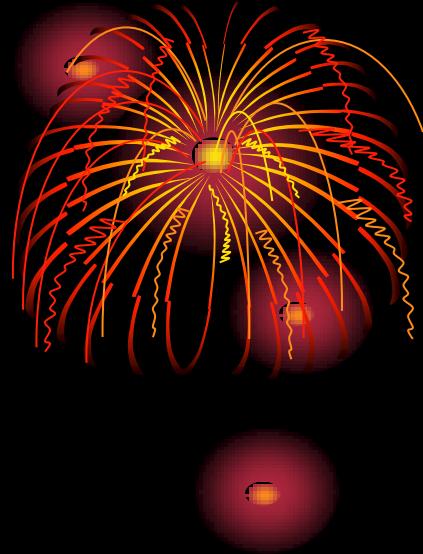
Sequential Quicksort



17	14	65	4	22	63	11
----	----	----	---	----	----	----

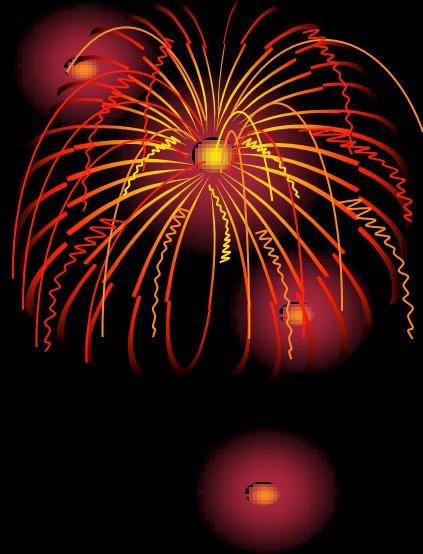
Unordered list of values

Sequential Quicksort



Choose pivot value

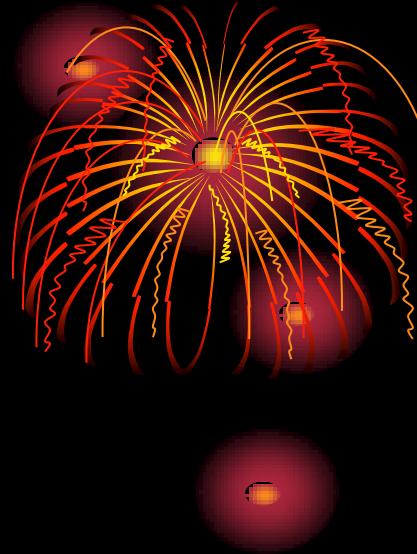
Sequential Quicksort



Low list
 (≤ 17)

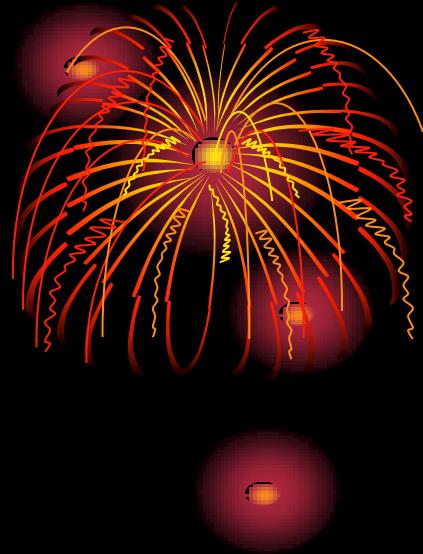
High list
 (> 17)

Sequential Quicksort



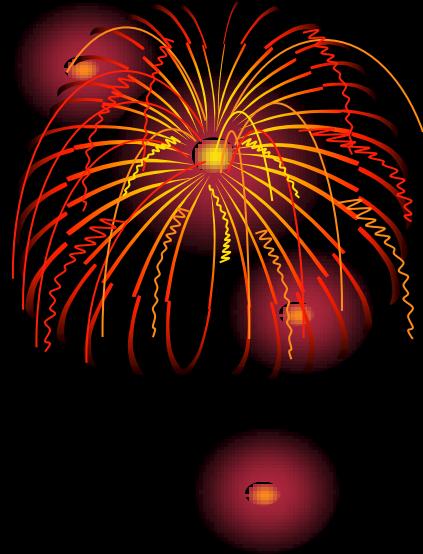
Recursively
apply quicksort
to low list

Sequential Quicksort



Recursively
apply quicksort
to high list

Sequential Quicksort



4	11	14	17	22	63	65
---	----	----	----	----	----	----

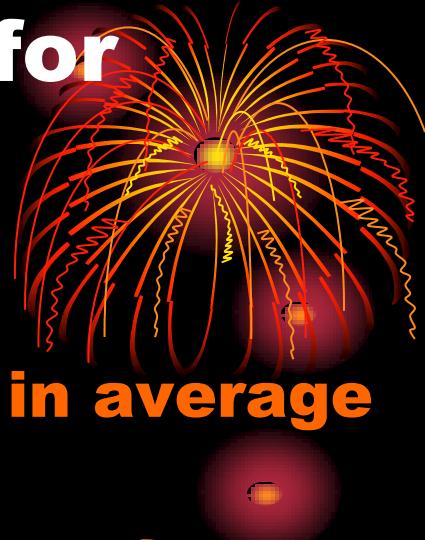
Sorted list of values

Attributes of Sequential Quicksort



- **Average-case time complexity:** $\Theta(n \log n)$
- **Worst-case time complexity:** $\Theta(n^2)$
 - **Occurs when low, high lists maximally unbalanced at every partitioning step**
- **Can make worst-case less probable by using sampling to choose pivot value**
 - **Example: “Median of 3” technique**

Quicksort Good Starting Point for Parallel Algorithm

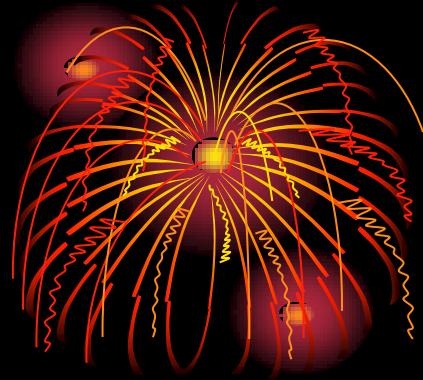


- **Speed**
 - **Generally recognized as fastest sort in average case**
 - **Preferable to base parallel algorithm on fastest sequential algorithm**
- **Natural concurrency**
 - **Recursive sorts of low, high lists can be done in parallel**

Definitions of “Sorted”

- **Definition 1: Sorted list held in memory of a single processor**
- **Definition 2:**
 - **Portion of list in every processor's memory is sorted**
 - **Value of last element on P_i 's list is less than or equal to value of first element on P_{i+1} 's list**
- **We adopt Definition 2: Allows problem size to scale with number of processors**

Parallel Quicksort



75, 91, 15, 64, 21, 8, 88, 54

P_0

50, 12, 47, 72, 65, 54, 66, 22

P_1

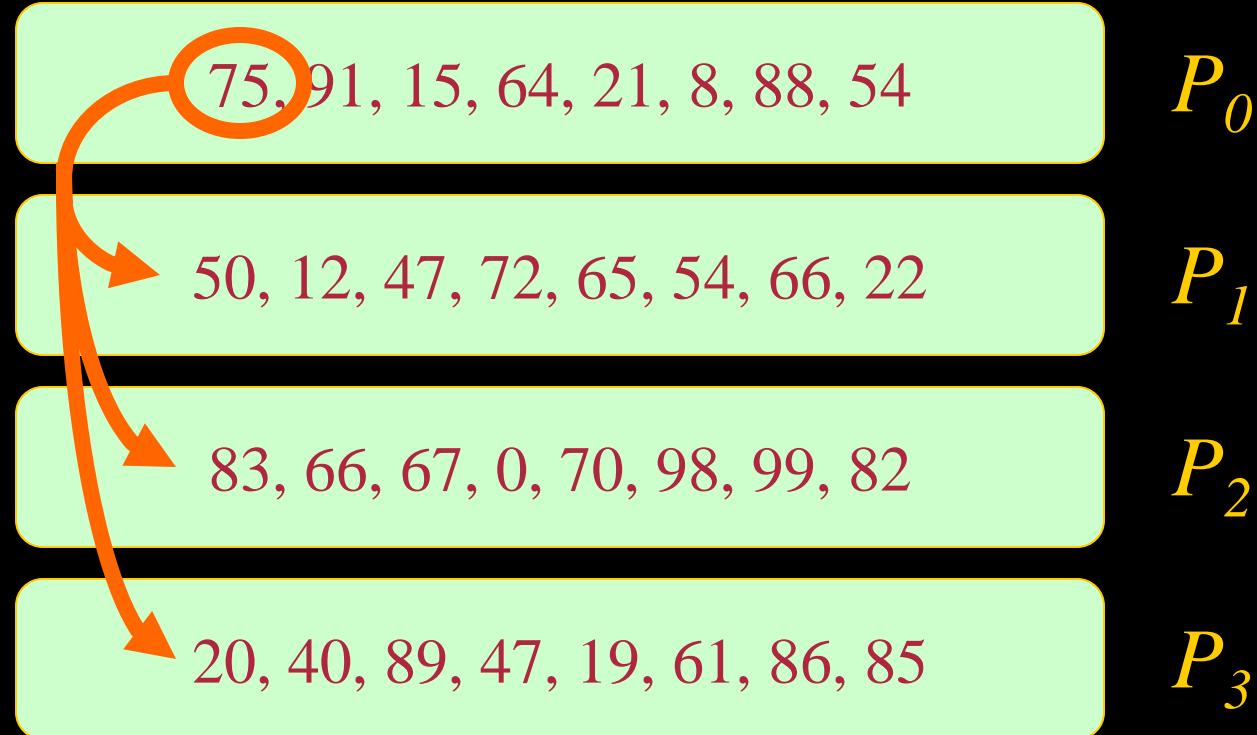
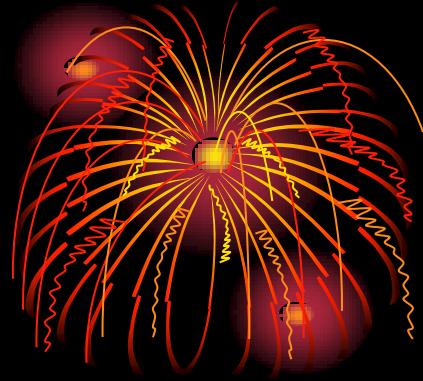
83, 66, 67, 0, 70, 98, 99, 82

P_2

20, 40, 89, 47, 19, 61, 86, 85

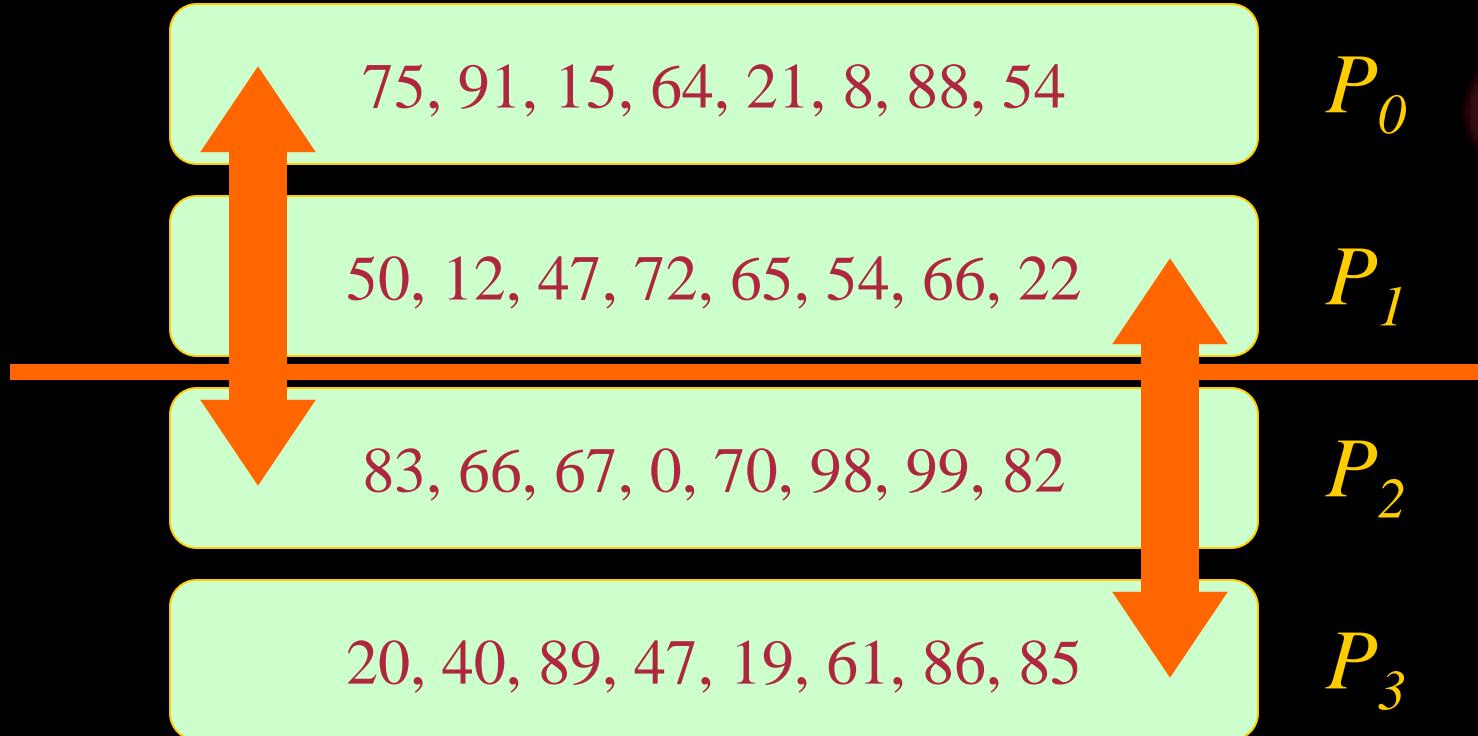
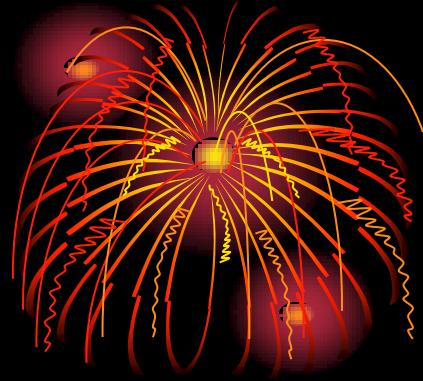
P_3

Parallel Quicksort



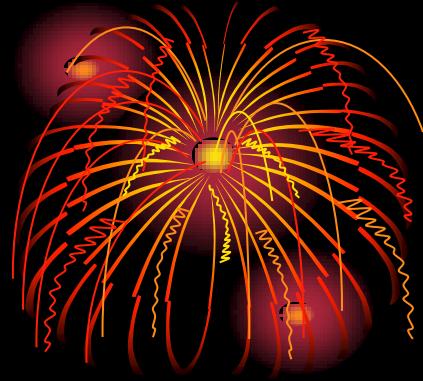
**Process P_0 chooses and broadcasts
randomly chosen pivot value**

Parallel Quicksort



Exchange “lower half” and “upper half” values

Parallel Quicksort



Lower
“half”

75, 15, 64, 21, 8, 54, 66, 67, 0, 70

P_0

Upper
“half”

50, 12, 47, 72, 65, 54, 66,
22, 20, 40, 47, 19, 61

P_1

83, 98, 99, 82, 91, 88

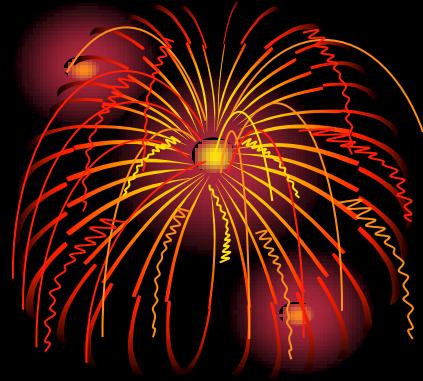
P_2

89, 86, 85

P_3

After exchange step

Parallel Quicksort



Lower
“half”

75, 15, 64, 21, 8, 54, 66, 67, 0, 70

P_0

50, 12, 47, 72, 65, 54, 66,
22, 20, 40, 47, 19, 61

P_1

Upper
“half”

83, 98, 99, 82, 91, 88

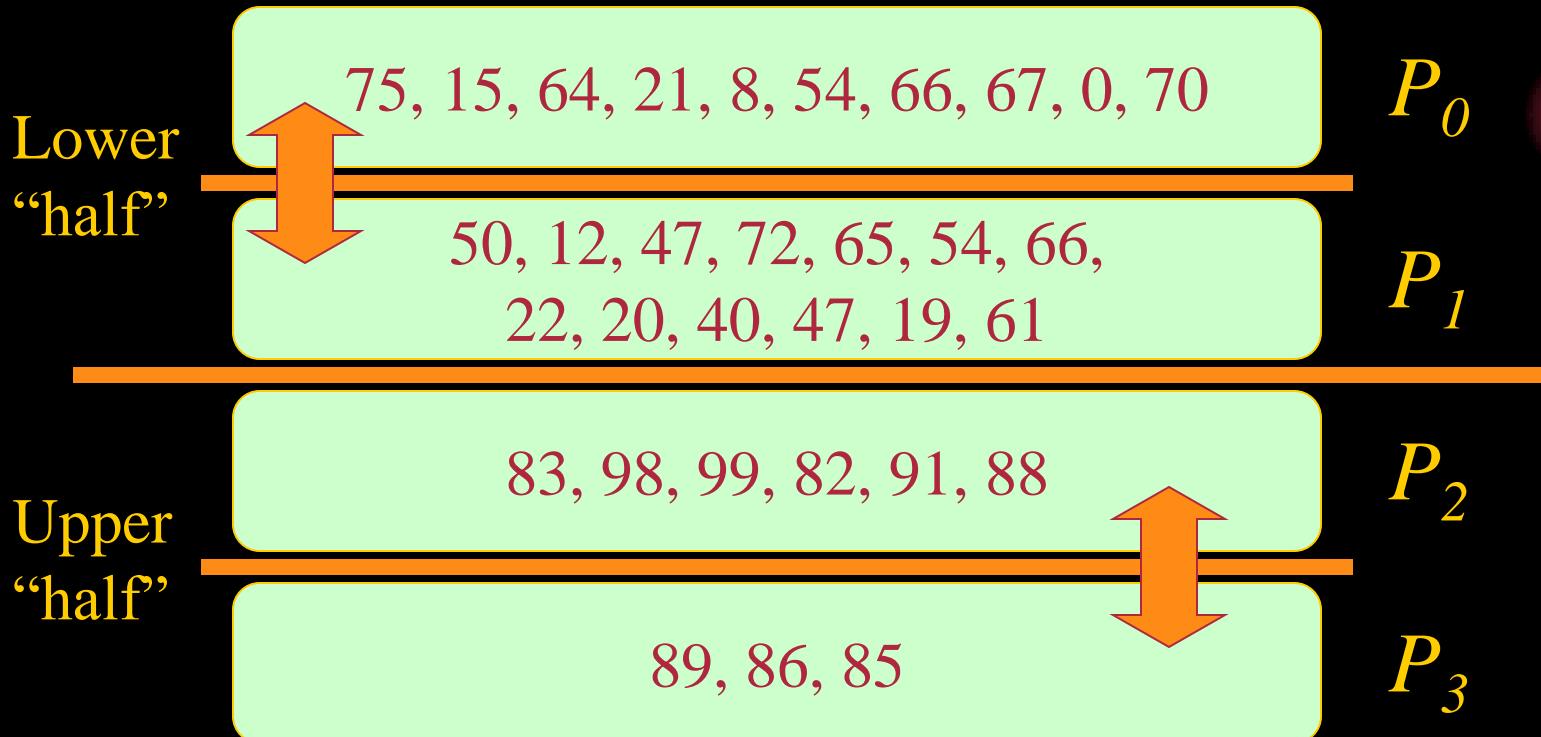
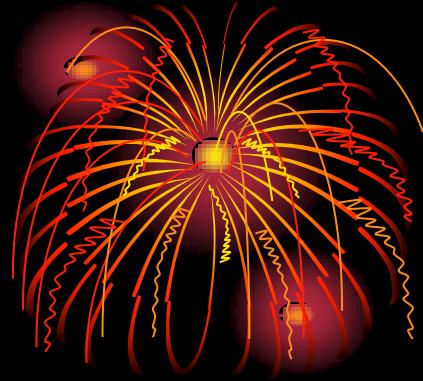
P_2

89, 86, 85

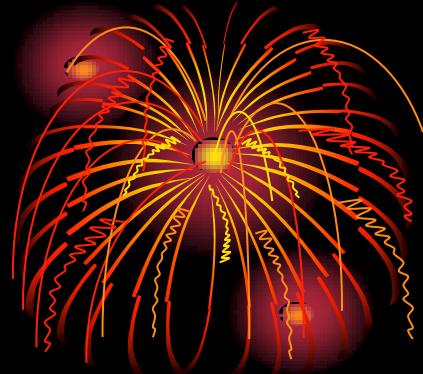
P_3

Processes P_0 and P_2 choose and broadcast randomly chosen pivots

Parallel Quicksort



Parallel Quicksort



Lower “half”
of lower “half”

15, 21, 8, 0, 12, 20, 19

P_0

Upper “half”
of lower “half”

50, 47, 72, 65, 54, 66, 22, 40,
47, 61, 75, 64, 54, 66, 67, 70

P_1

Lower “half”
of upper “half”

83, 82, 91, 88, 89, 86, 85

P_2

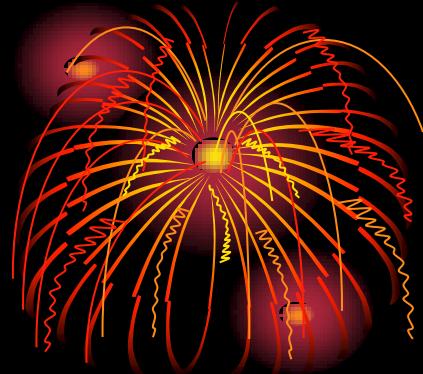
Upper “half”
of upper “half”

98, 99

P_3

After exchanging values

Parallel Quicksort



Lower “half”
of lower “half”

0, 8, 12, 15, 19, 20, 21

P_0

Upper “half”
of lower “half”

22, 40, 47, 47, 50, 54, 54, 61,
64, 65, 66, 66, 67, 70, 72, 75

P_1

Lower “half”
of upper “half”

82, 83, 85, 86, 88, 89, 91

P_2

Upper “half”
of upper “half”

98, 99

P_3

Each processor sorts values it controls

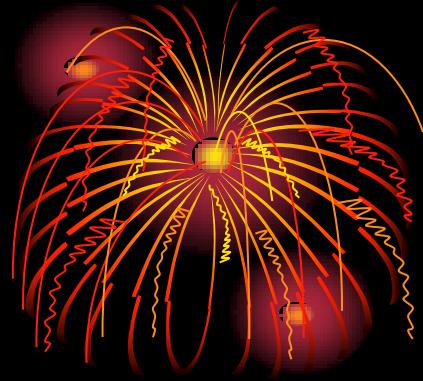
Analysis of Parallel Quicksort

- **Execution time dictated by when last process completes**
- **Algorithm likely to do a poor job balancing number of elements sorted by each process**
- **Cannot expect pivot value to be true median**
- **Can choose a better pivot value**

Hyperquicksort

- **Start where parallel quicksort ends: each process sorts its sublist**
- **First “sortedness” condition is met**
- **To meet second, processes must still exchange values**
- **Process can use median of its sorted list as the pivot value**
- **This is much more likely to be close to the true median**

Hyperquicksort



75, 91, 15, 64, 21, 8, 88, 54

P_0

50, 12, 47, 72, 65, 54, 66, 22

P_1

83, 66, 67, 0, 70, 98, 99, 82

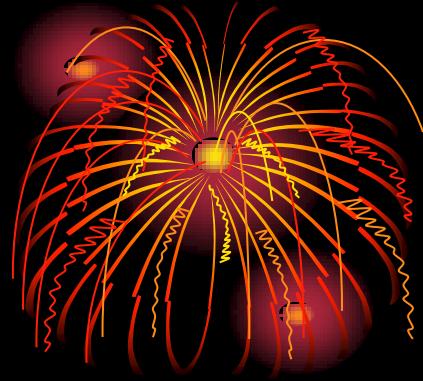
P_2

20, 40, 89, 47, 19, 61, 86, 85

P_3

Number of processors is a power of 2

Hyperquicksort



8, 15, 21, 54, 64, 75, 88, 91

P_0

12, 22, 47, 50, 54, 65, 66, 72

P_1

0, 66, 67, 70, 82, 83, 98, 99

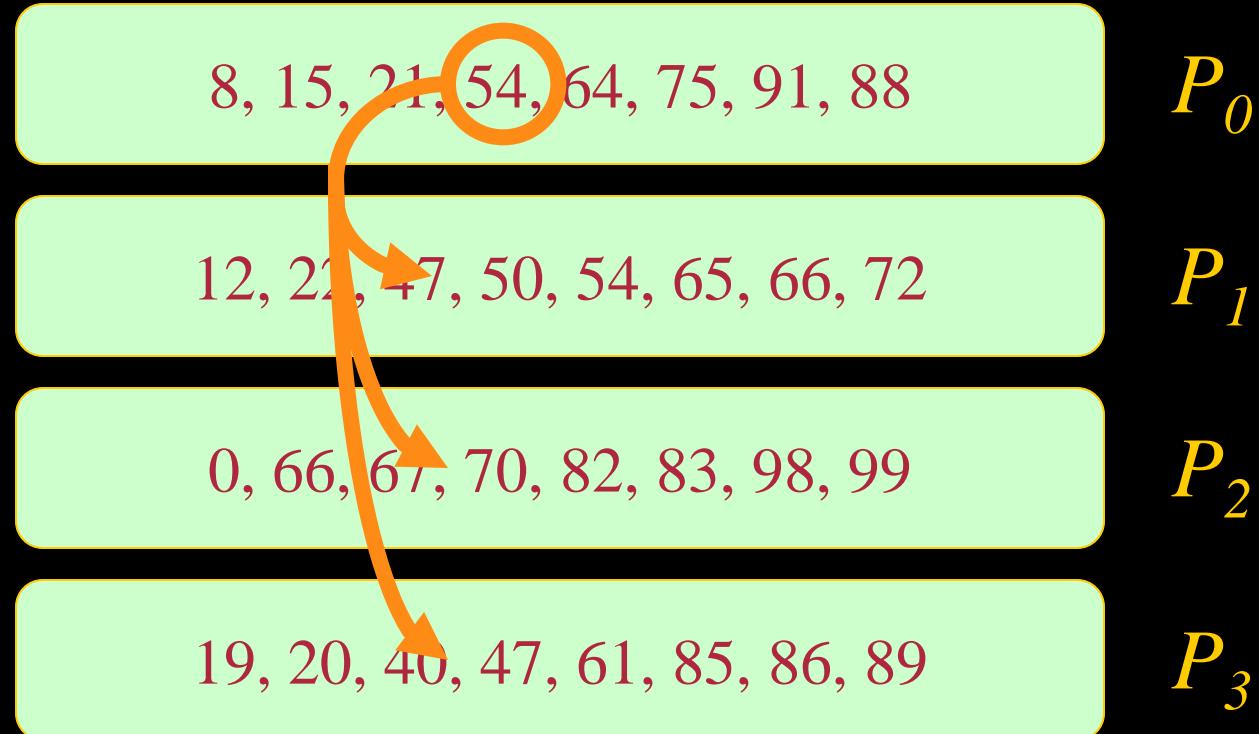
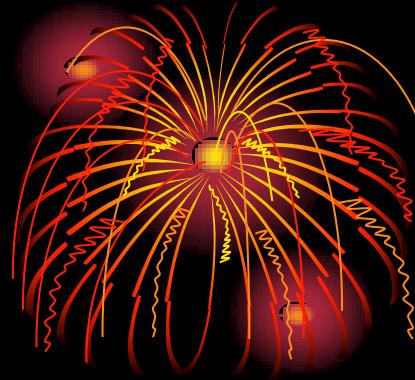
P_2

19, 20, 40, 47, 61, 85, 86, 89

P_3

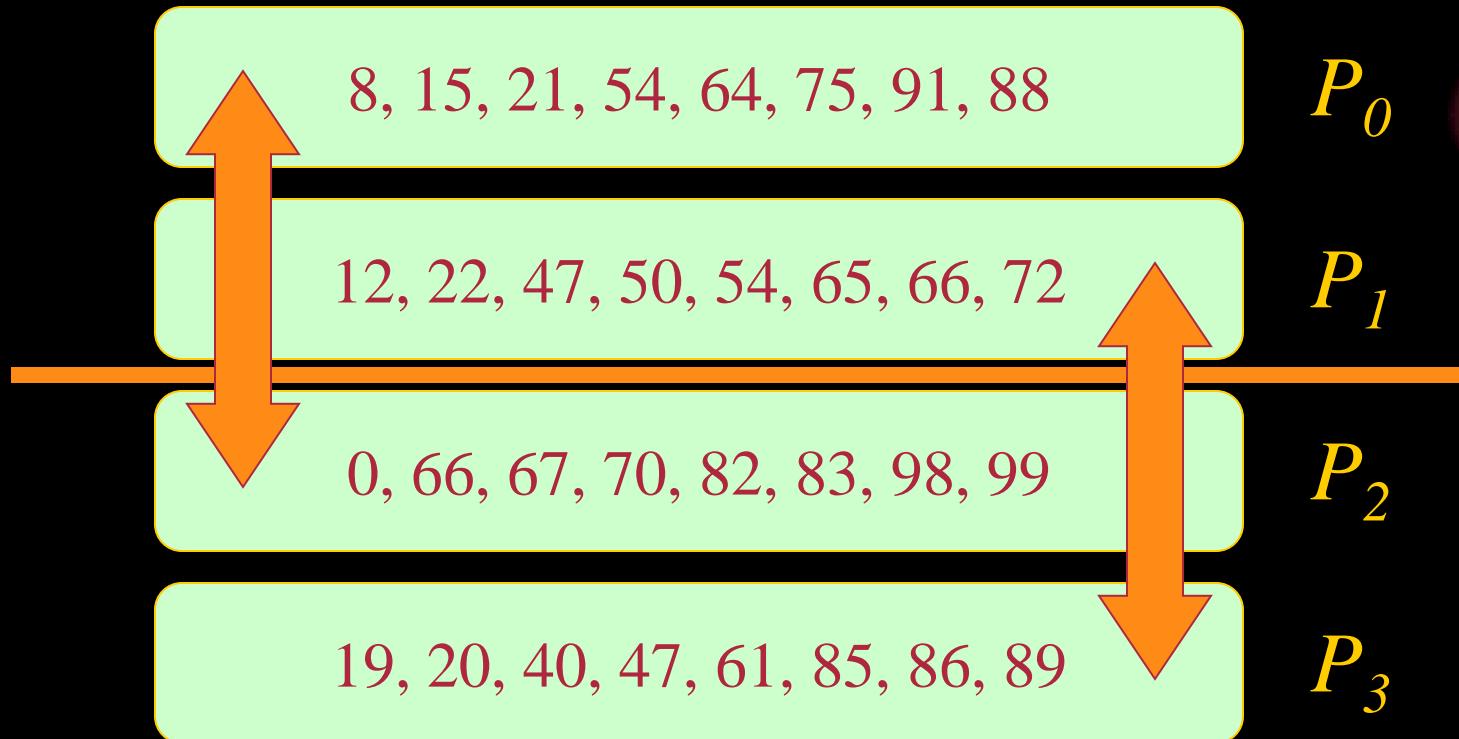
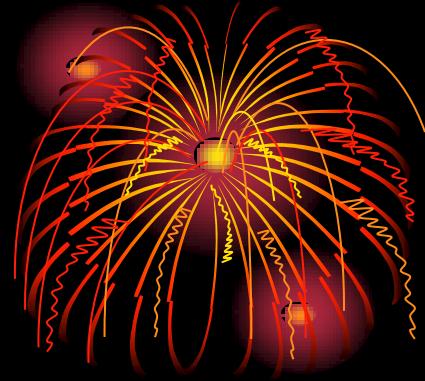
Each process sorts values it controls

Hyperquicksort



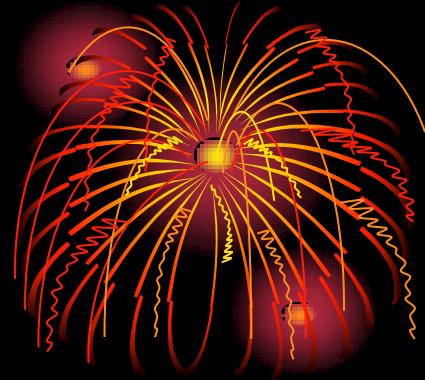
Process P_0 broadcasts its median value

Hyperquicksort



Processes will exchange “low”, “high” lists

Hyperquicksort



0, 8, 15, 21, 54

P_0

12, 19, 20, 22, 40, 47, 47, 50, 54

P_1

64, 66, 67, 70, 75, 82, 83, 88, 91, 98, 99

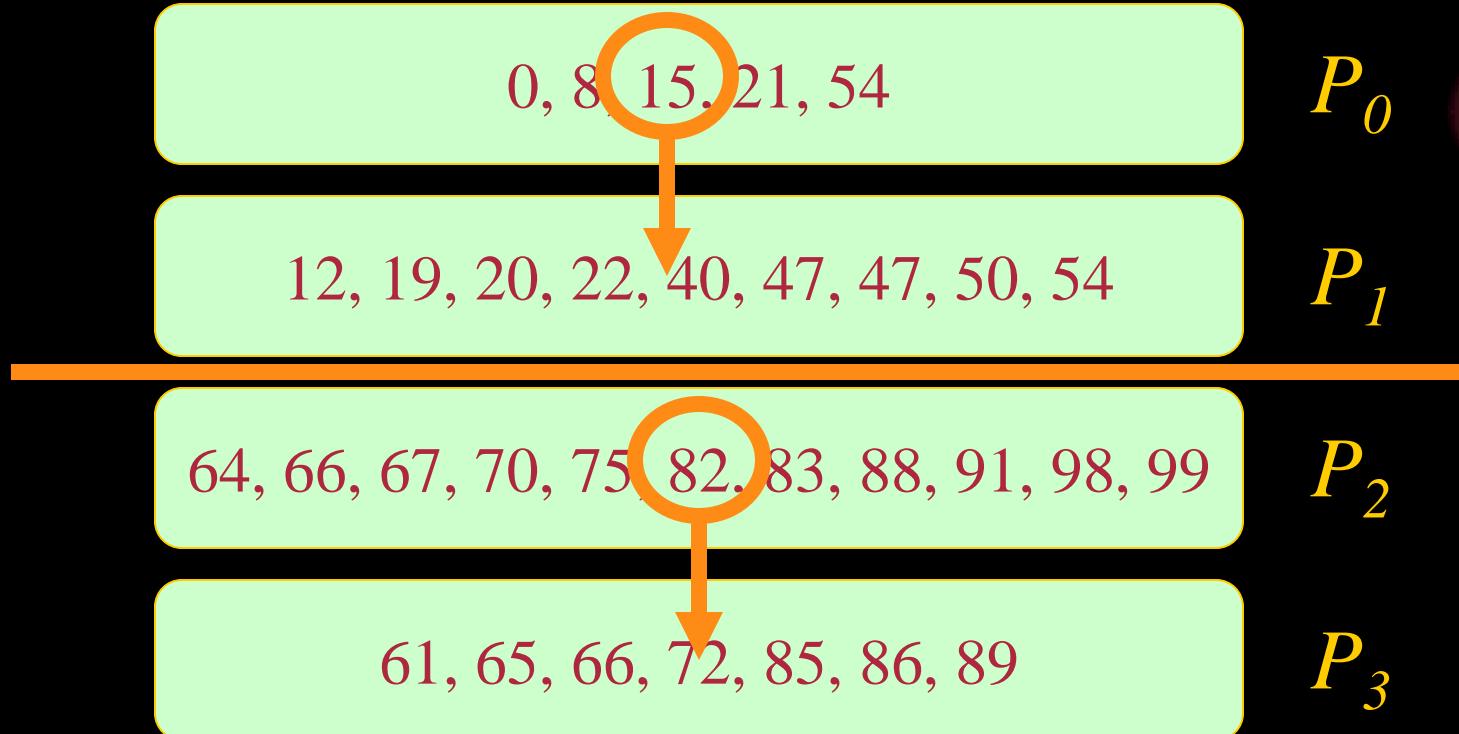
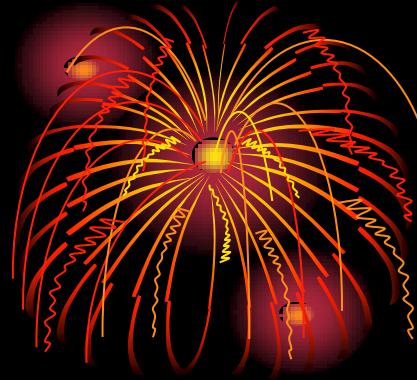
P_2

61, 65, 66, 72, 85, 86, 89

P_3

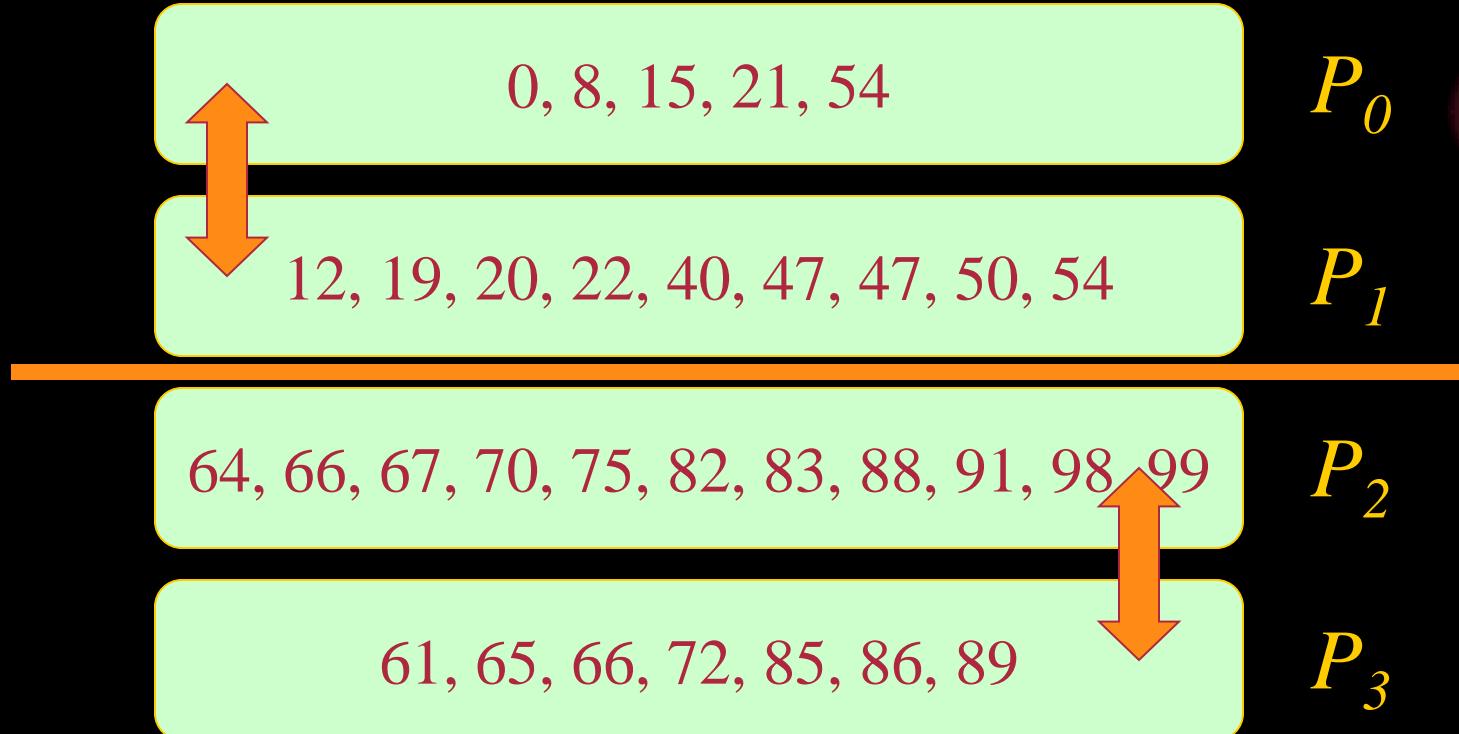
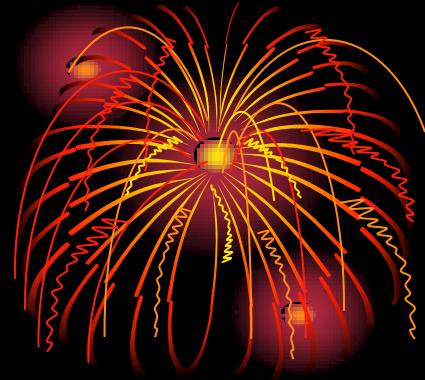
Processes merge kept and received values.

Hyperquicksort



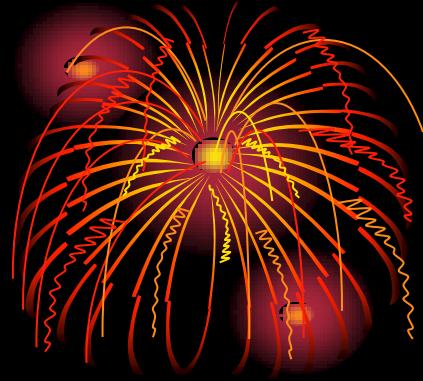
Processes P_0 and P_2 broadcast median values.

Hyperquicksort



Communication pattern for second exchange

Hyperquicksort



0, 8, 12, 15

P_0

19, 20, 21, 22, 40, 47, 47, 50, 54, 54

P_1

61, 64, 65, 66, 66, 67, 70, 72, 75, 82

P_2

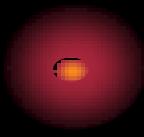
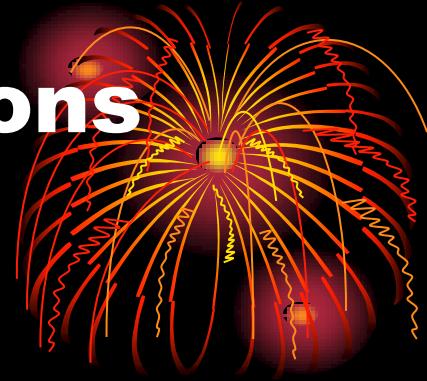
83, 85, 86, 88, 89, 91, 98, 99

P_3

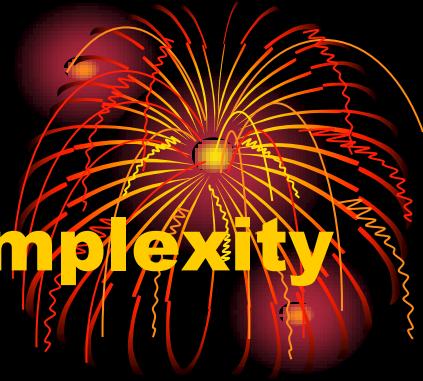
After exchange-and-merge step

Complexity Analysis Assumptions

- **Average-case analysis**
- **Lists stay reasonably balanced**
- **Communication time dominated by message transmission time, rather than message latency**



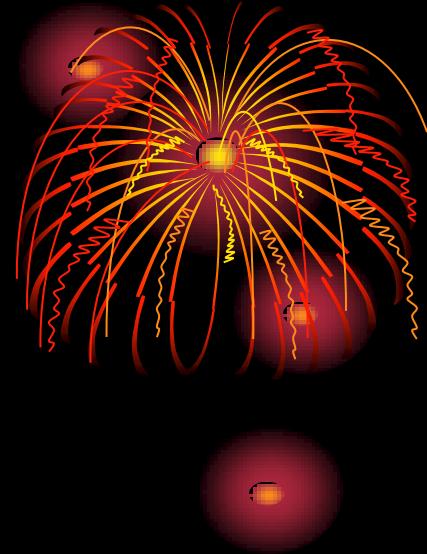
Complexity Analysis



- **Initial quicksort step has time complexity $\Theta((n/p) \log (n/p))$**
- **Total comparisons needed for $\log p$ merge steps: $\Theta((n/p) \log p)$**
- **Total communication time for $\log p$ exchange steps: $\Theta((n/p) \log p)$**

Isoefficiency Analysis

- Sequential time complexity: $\Theta(n \log n)$
- Parallel overhead: $\Theta(n \log p)$
- Isoefficiency relation:
 $n \log n \geq C n \log p \Rightarrow \log n \geq C \log p \Rightarrow n \geq p^C$



$$M(p^C)/p = p^C / p = p^{C-1}$$

- The value of C determines the scalability. Scalability depends on ratio of communication speed to computation speed.

Another Scalability Concern

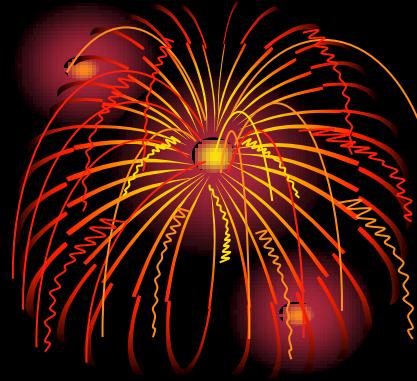
- Our analysis assumes lists remain balanced
- As p increases, each processor's share of list decreases
- Hence as p increases, likelihood of lists becoming unbalanced increases
- Unbalanced lists lower efficiency
- Would be better to get sample values from all processes before choosing median

Parallel Sorting by Regular Sampling (PSRS Algorithm)



- **Each process sorts its share of elements**
- **Each process selects regular sample of sorted list**
- **One process gathers and sorts samples, chooses pivot values from sorted sample list, and broadcasts these pivot values**
- **Each process partitions its list into p pieces, using pivot values**
- **Each process sends partitions to other processes**
- **Each process merges its partitions**

PSRS Algorithm



75, 91, 15, 64, 21, 8, 88, 54

P_0

50, 12, 47, 72, 65, 54, 66, 22

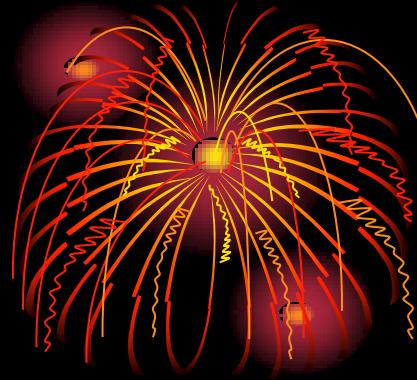
P_1

83, 66, 67, 0, 70, 98, 99, 82

P_2

Number of processors does not have to be a power of 2.

PSRS Algorithm



8, 15, 21, 54, 64, 75, 88, 91

P_0

12, 22, 47, 50, 54, 65, 66, 72

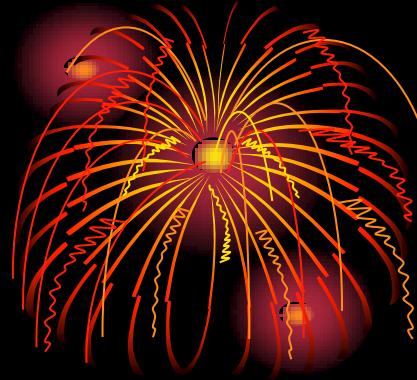
P_1

0, 66, 67, 70, 82, 83, 98, 99

P_2

Each process sorts its list using quicksort.

PSRS Algorithm



P_0

8, 15 21, 54 64, 75 88, 91

P_1

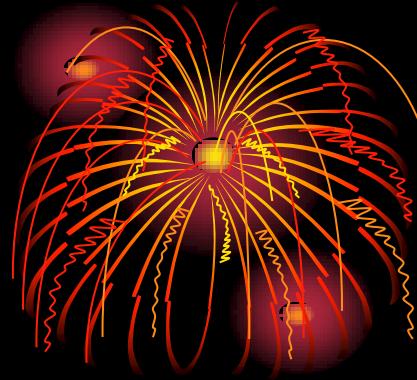
12, 22 47, 50 54, 65 66, 72

P_2

0, 66 67, 70 82, 83 98, 99

Each process chooses p regular samples.

PSRS Algorithm



P_0

8, 15 21, 54 64, 75 88, 91

P_1

12, 22 47, 50 54, 65 66, 72

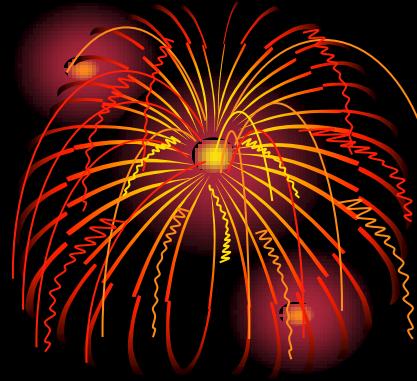
P_2

0, 66 67, 70 82, 83 98, 99

15, 54, 75, 22, 50, 65, 66, 70, 83

One process collects p^2 regular samples.

PSRS Algorithm



P_0

8, 15 21, 54 64, 75 88, 91

P_1

12, 22 47, 50 54, 65 66, 72

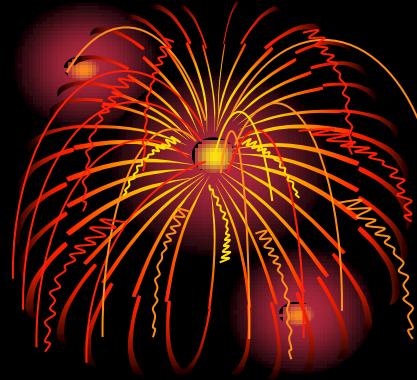
P_2

0, 66 67, 70 82, 83 98, 99

15, 22, 50, 54, 65, 66, 70, 75, 83

The same process sorts the p^2 regular samples.

PSRS Algorithm



8, 15, 21, 54, 64, 75, 88, 91

P_0

12, 22, 47, 50, 54, 65, 66, 72

P_1

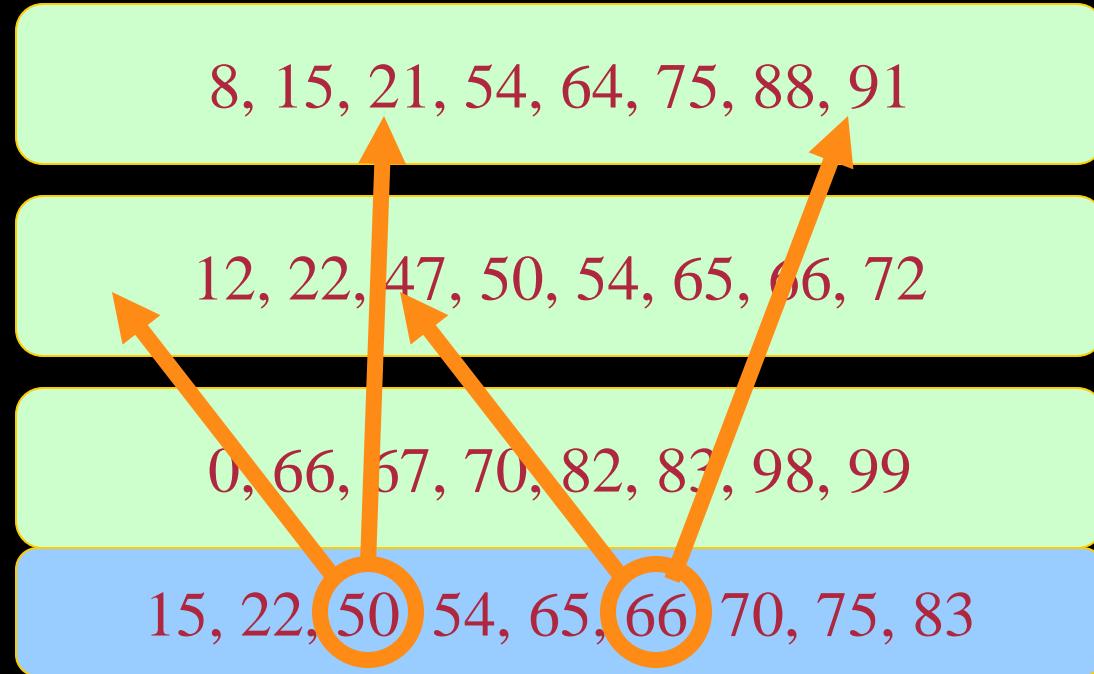
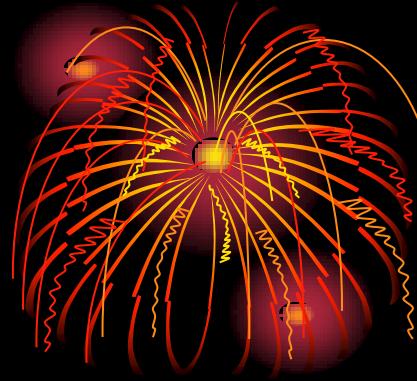
0, 66, 67, 70, 82, 83, 98, 99

P_2

15, 22, 50 54, 65, 66 70, 75, 83

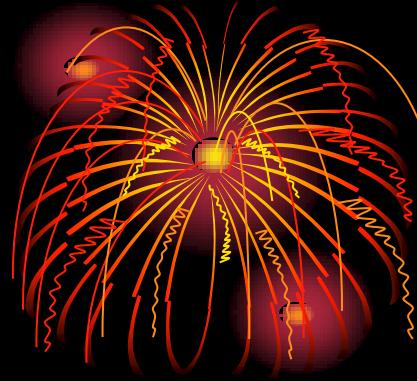
The process then chooses $p-1$ pivot values.

PSRS Algorithm



The process broadcasts the $p-1$ pivot values.

PSRS Algorithm



P_0

8, 15, 21, 54, 64, 75, 88, 91

P_1

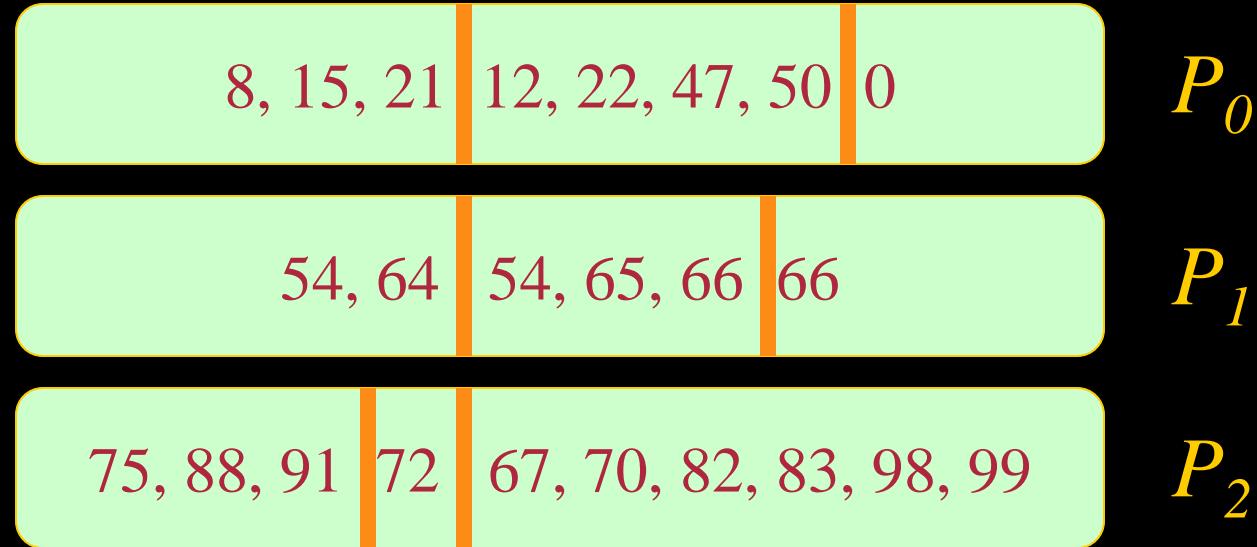
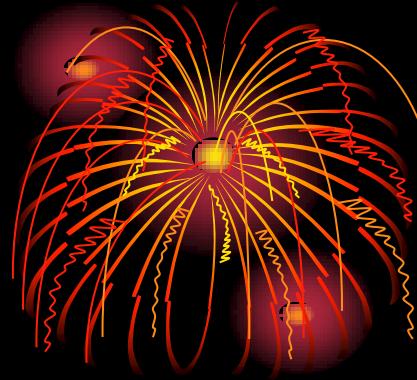
12, 22, 47, 50, 54, 65, 66, 72

P_2

0, 66, 67, 70, 82, 83, 98, 99

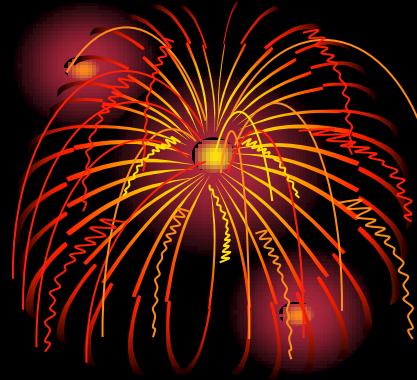
Each process divides its list, based on the pivot values.

PSRS Algorithm



Each process sends its partitions to the correct destination process.

PSRS Algorithm



0, 8, 12, 15, 21, 22, 47, 50

P_0

54, 54, 64, 65, 66, 66

P_1

67, 70, 72, 75, 82, 83, 88, 91, 98, 99

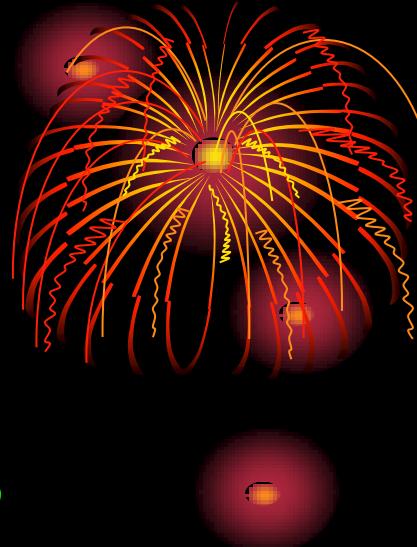
P_2

Each process merges its p partitions.

Assumptions

- **Each process ends up merging close to n/p elements**
- **Experimental results show this is a valid assumption**
- **Processor interconnection network supports p simultaneous message transmissions at full speed (full-duplex full switch)**

Time Complexity Analysis



- **Computations**

- **Initial quicksort:** $\Theta((n/p)\log(n/p))$
- **Sorting regular samples:** $\Theta(p^2 \log p)$
- **Merging sorted sublists:** $\Theta((n/p)\log p)$
- **Overall:** $\Theta((n/p)(\log n + \log p) + p^2 \log p)$

- **Communications**

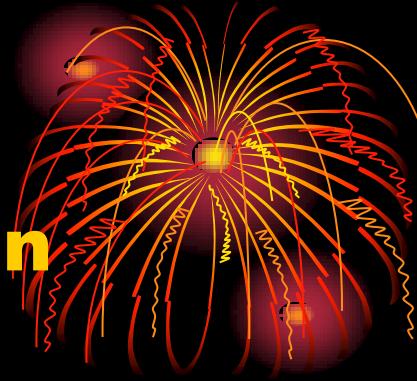
- **Gather samples, broadcast pivots:** $\Theta(\log p)$
- **All-to-all exchange:** $\Theta(n/p)$
- **Overall:** $\Theta(n/p)$

Isoefficiency Analysis



- Sequential time complexity: $\Theta(n \log n)$
- Parallel overhead: $\Theta(n \log p)$
- Isoefficiency relation:
 $n \log n \geq Cn \log p \Rightarrow \log n \geq C \log p$
- Scalability function same as for hyperquicksort
- Scalability depends on ratio of communication to computation speeds

Summary



- **Three parallel algorithms based on quicksort**
- **Keeping list sizes balanced**
 - Parallel quicksort: poor
 - Hyperquicksort: better
 - PSRS algorithm: excellent
- **Average number of times each key moved:**
 - Parallel quicksort and hyperquicksort: $\log p / 2$
 - PSRS algorithm: $(p-1)/p$